

AP-BidIQ — Pitch Deck

Slot: 15 min + Q&A · **Audience:** Mixed govt + technical · **Format:** Remote (Zoom/Teams)

Slide 1 — Cover

AP-BidIQ

AI-Powered Bid Evaluation for AP Procurement

Submitted to: AP Hackathon — *AI for Governance and Procurement Efficiency* Department: Infrastructure & Investment, Government of Andhra Pradesh

Cittaai · 2026-04-30

Slide 2 — The Problem

One tender. 16 sections. 19 forms. 1 corrigendum. 1 officer. 1 week.

The EPCC Fishing Harbours Phase II package you sent us: - 16 documents totalling **~1 MB of structured legal text** - **19 vendor forms** to verify per bid - **15 mandatory clauses** to check - **Quantitative thresholds** scattered across TDS, ITT, QR, and a corrigendum - **Corrigendum-1** silently changed the bid-capacity lookback from 5 to 10 years


Today an officer evaluates 5 bids per tender, manually, in **2-3 days each**. Errors trigger appeals. Delays compound. Evaluators carry the audit risk.

Slide 3 — The Cost of the Status Quo

Pain	Impact
Manual checklist work	70% of officer time on mechanical reconciliation
Inconsistent application of corrigenda	Disputed verdicts, appeals, re-tenders
Subjective evaluation	Vendor mistrust, RTI queries, judicial intervention
No audit trail of <i>why</i> a bid was rejected	Hard-to-defend decisions; officer personal risk

Slide 4 — Solution

Versioned Ingest → Layered Validation → Defensible Citation → Officer Decides



AI never auto-rejects. Officer reviews every finding with a 1-click source citation.

Slide 5 — LIVE DEMO

Switch to browser — open the deployed app.

Demo flow (~7 min, mirrors §3 of DEMO_PLAN): 1. Tender page → Corrigendum-1 diff (5→10 yr lookback) 2. Bid list → 5 vendors evaluated 3. Click bid → 13-check scorecard 4. Click finding → citation drawer (quoted source) 5. Bid capacity formula failure: $(45 \times 3 \times 3) - 280 = ₹125 \text{ Cr} < ₹380 \text{ Cr}$ 6. Generated Evaluation Statement (matches AP human format) 7. Comparison view → 1 of 5 qualifies for L1 8. Audit log → "✓ Chain verified · 25 entries"

Slide 6 — What We Hit on Your Corpus

Brief targets — exceeded.

Metric	Brief target	
Defect detection accuracy	≥ 85%	100.0%
Verdict match with human	≥ 90%	100.0%
False positive count	(lower better)	0

Tested on **5 vendor bids** (1 clean baseline + 4 with planted defects mirroring real audit findings). 13 compliance checks per bid (10 deterministic + 3 LLM-driven).

Caveat surfaced proactively: small N. Production rollout extends benchmarking to the 50+ historical bids per the brief.

Slide 7 — Architecture

Layer	Stack	Why
Frontend	Next.js 16 + Tailwind	Officer dashboard; sub-second navigation
API	FastAPI (Python 3.12)	Best ecosystem for doc parsing + LLM SDKs
Database	Postgres (Neon) + SQLite for dev	Transparent dev/prod parity via SQLAlchemy
AI	OpenAI GPT-4o (clause semantics + drafting) + GPT-4o-mini (batch extraction)	Provider-flexible architecture — Claude / Llama / Mistral are single-file swaps. Built-in prompt caching cuts cost on repeated tender context.
Audit	SHA-256 hash chain in Postgres	Tamper-evident; verifies in 1 SQL query
Hosting	Vercel + Railway + Neon	All free tier for the demo; on-prem option available

Cost per bid evaluated: ~₹2 (3 LLM calls × cached context). **Latency:** < 60s end-to-end for a fresh upload.

Six brief features, all built: Drafting Assistant · Validation · Bid Evaluation · Document Comparison (cross-bid anomaly) · Officer Dashboard · Audit Trail. Plus PII redaction, English+Telugu UI, and OCR for scanned bids.

Slide 8 — Why Us, Why Now

Differentiator	Most generic AI procurement tools	
Built on AP's actual tender	✗ — retrofit demos	✓ — your EPCC + Corrigendum-1
Corrigendum versioning	✗ — ignored	✓ — versioned ActiveRules with diff view
Officer-defensibility	✗ — black-box scores	✓ — every finding cites quoted source
Bid-capacity formula compute	✗ — only NLP	✓ — $A \times N \times 3 - B$ deterministic, with corrigendum-aware $A_lookback$
Tamper-evident audit	✗ — log files	✓ — hash chain, in-app verifier

Bottom line: *We didn't bring a generic tool. We brought your tool.*

Slide 9 — Roadmap (post-pilot)

Quarter	Module
Q1	Drafting Assistant — RAG over AP clause library + project-type templates
Q2	Multi-department generalisation — works, services, goods, consultancy
Q3	Vendor portal + e-procurement portal API integration
Q4	Cross-bid anomaly detection — collusion/cartelisation patterns

Beyond Q4: Telugu UI · OCR for scanned bids · Mobile officer app · API for state CAG audit access.

Slide 10 — The Ask

30-day pilot in the Infrastructure & Investment Department

What we'd do: - Onboard **50 historical bids** to validate accuracy in real conditions - Add the **drafting assistant** (cut from this demo for time) - Add **bilingual support** if any of those bids include Telugu content - White-glove training for **2 procurement officers** - All AP data stays in your AWS region; full export at end of pilot

Cost to AP for the 30-day pilot: ₹0 (we absorb compute + labour).

Slide 11 — Contact

Mohineesh / Cittaai ✉ tech@cittaai.com 🔗 GitHub: github.com/cma-org/ap-bidiq (*repo link*) 🌐 Live demo: ap-bidiq.vercel.app (*deployed URL*)

Leave-behind kit: 2-page brief + this deck + 90-sec demo video — all in your inbox.

Appendix

A — Cost economics

3 Sonnet calls × 600 input + 150 output tokens, 90% prompt-cache hit on shared RFP context = **₹2 / bid evaluated**.
At 1,000 bids/month = ₹2,000.

B — Privacy & data residency

- AP data never used for model training (Claude API contract).
- Optional self-hosted: Llama 3.1 70B on AP infrastructure (~3-week migration).
- PII redaction at ingest (vendor contact info masked before LLM call).

C — Anticipated questions

(See Q&A section of DEMO_PLAN.md — 20 questions pre-answered.)